

TEJAS PANDYA

Brooklyn, NY | tbp8777@nyu.edu | (551) 376-9633

LinkedIn: [linkedin.com/in/tejaspandya9598](https://www.linkedin.com/in/tejaspandya9598) | GitHub: github.com/tejaspandya9598 | Portfolio: tejaspandya.me

EDUCATION

NEW YORK UNIVERSITY, TANDON SCHOOL OF ENGINEERING , Brooklyn, NY	Expected 05/27
Master of Science in Financial Engineering GPA: 3.8/4.0	
Graduate Teaching Assistant: Financial Economics; C++ Financial Software Laboratory; Value Investing	
SVKM'S NMIMS, SCHOOL OF BUSINESS MANAGEMENT , Mumbai, India	04/22
Master of Business Administration, Concentration in Finance GPA: 3.3/4.0	
VELLORE INSTITUTE OF TECHNOLOGY , Vellore, India	06/20
Bachelor of Technology in Mechanical Engineering GPA: 8.9/10.0	

PROGRAMMING / TECHNICAL SKILLS / CERTIFICATIONS

Programming: C/C++ (C++20, STL, Boost), Python, SQL, Bash, LaTeX/Typst
Libraries & Frameworks: NumPy, Pandas, Polars, SciPy, scikit-learn, statsmodels, LightGBM, XGBoost, PyTorch, pytest
Data & Tools: Linux/Unix, Git, CMake, Streamlit, Shiny, Excel, Bloomberg Terminal, Tableau, Power BI
Certifications: CFA Level I; C++ Programming for Financial Engineering (Baruch College); Bloomberg Market Concepts; CFI FMVA

COURSEWORK HIGHLIGHTS

Quantitative Methods & Computing: Financial Econometrics, Machine Learning in Financial Engineering, Quantitative Methods in Finance, Financial Computing, Applied Numerical Methods, Data Structures and Algorithms
Risk, Valuation & Derivatives: Financial Risk Management, Valuation for Financial Engineering, Quantitative Portfolio Management, Options Pricing & Stochastic Calculus, Fixed Income Securities and Interest Rate Derivatives

EXPERIENCE

NEW YORK UNIVERSITY, TANDON SCHOOL OF ENGINEERING , Brooklyn, NY	05/26 - 08/26
<i>Quantitative Researcher, QWIM Industry Capstone</i>	
- Lifted Sharpe from 0.37 to 0.55 against a 60/40 benchmark over 1999-2025 as one of five on the QWIM capstone, pairing equal-risk-contribution weighting with hierarchical budgeting on Conditional Drawdown-at-Risk and blending regime-specific covariance by one-month-ahead probability - Compressed annualised volatility from 11.5% to 6.3%, maximum drawdown from 37% to 24% and Conditional Drawdown-at-Risk from 0.250 to 0.169 against the same 60/40 benchmark, accepting a marginally lower absolute return of 5.4% against 5.7% and turnover of 0.242 - Implemented a log-periodic power-law singularity bubble-confidence tilt as the fourth ablation level above regime-conditioned covariance blending and quantile-regression tail estimation, the strongest of eight variants, holding Sharpe 0.56 against 0.50 on the 2008-2025 window	
NOMURA SERVICES INDIA PRIVATE LIMITED , Mumbai, India	08/22 - 07/25
<i>Senior Risk Analyst, Market Risk Infrastructure</i>	
- Owned the lifecycle of VaR and sensitivity models across equities, rates, FX and commodities, from daily production runs through backtesting; Python anomaly detection triaged exceptions to stale marks, erroneous data or booking errors, cutting investigation turnaround 40% and restoring SLAs - Delivered quarterly Basel 2.5 and FRTB stress-testing and capital submissions alongside front-office and regulatory teams, engineering rules-based validation over delta, vega and curvature sensitivities that reconciled fully against front-office pricing and caught anomalies pre-submission - Engineered Python and SQL pipelines across equities, rates, FX and commodities that aggregate, validate and reconcile daily P&L, positions and Level-1/Level-2 market data into the VaR and sensitivity dashboards, eliminating more than 50 hours of manual processing every month - Led a three-person workstream on US-region end-of-day risk reporting, cutting overnight batch runtime 35% through SQL query optimisation and job scheduling, and built the VaR, P&L and exception dashboards that senior stakeholders read each morning before the US open	

ACADEMIC PROJECTS

TREASURY-CURVE SWAP VALUATION & HISTORICAL COUNTERPARTY RISK	05/26
<i>Python, NumPy, SciPy</i> github.com/tejaspandya9598/sofr-swap-pnl-attribution	
- Bootstrapped discount curves from FRED Treasury constant-maturity par yields rather than licensed SOFR fixings, marking a \$255M ten-swap 2-30 year book and decomposing daily P&L into carry, roll-down and orthogonal level, slope and curvature moves reconciling to the total within 1e-6 - Attributed a 51-76bp one-day curve collapse on 16 April 1980 to carry, roll-down and the orthogonal level, slope and curvature factors, booking the 9% convexity residual explicitly against 0.9% on a quiet day, the honest limit of a first-order DV01 decomposition rather than a hidden plug	
ALMGREN-CHRISS OPTIMAL EXECUTION	05/26
<i>Python, NumPy, Deribit API</i> github.com/tejaspandya9598/optimal-execution	
- Trained a NumPy policy-gradient agent with hand-derived gradients checked against central differences in continuous integration, on 207,361 real five-minute bars; it converged to within 0.060bps of the closed form it was never shown, and nothing beat TWAP out of sample - Built the episode simulator over 5,752 non-overlapping three-hour episodes split chronologically seventy-thirty, with volatility, volume and a Corwin-Schultz half-spread estimated only from the 288 bars closing before each episode opened, verified by a test that triples every in-window bar	
HFT LIMIT ORDER BOOK & MATCHING ENGINE	02/26
<i>C++20, CMake, CTest</i> github.com/tejaspandya9598/hft-matching-engine	
- Built a deterministic price-time-priority matching engine timing roughly 21M messages per second on an Apple M2 over a pre-generated synthetic message mix, and 10.7M through the lock-free end-to-end pipeline, on a zero-allocation hot path with 64-byte cache-aligned structures - Reached O(1) average add, cancel and execute through intrusive per-level lists with hash lookup, a new price level costing O(log L) for best-bid tracking; 20 CTest assertions plus five benchmark checks pin price-time priority, partial fills and pool reuse under 10k-order churn	